

Exercise 2: LeNet in Tensorflow

David-Elias Künstle

November 20, 2017

1 Introduction

In this report I present two small experiments on handwritten digit detection (*MNIST dataset*) with a LeNet-like convolutional neural network. The network consists of two convolutional layers with 16 3x3 filters (stride 1, same padding), both with ReLu activation and max pooling, followed by a fully connected layer with 128 hidden units and 10 output units after a softmax. It is implemented in *python3* using the *Tensorflow library* (`tensorflow-gpu 1.3.0`). For training 50.000 digit-label pairs were used. The validation set contains 10.000 ones which are not in the training set.

2 Changing the Learning Rate

The influence of learning rate on validation error is what we see in Figure 1. You see the lower error rate while all epochs the higher the learning rate (α) is.

For $\alpha \in \{0.001, 0.01, 0.1\}$ error rate drops in the first 1–4 epochs and converges afterwards. Still the steepness while the drop and the value of the limit differ depending on α . I assume, that the error rates after training will stay almost constant, even if we would add some epochs. Instead the course for $\alpha = 0.0001$ is looks like linear decrease. It's highly likely that while additional training epochs the error rate could decrease further and start converging later.

I would prefer $\alpha = 0.1$ because it converges quickly to the lowest error rate. Training could be stopped after just one or two epochs. Nonetheless we know, that high learning rates increase the risk for overshooting and other problems while training.




Figure 1: Error rate on validation dataset of LeNet-like CNN while ongoing training with various learning rates.

3 Runtime on CPU vs GPU

Each value of the filters in a convolutional layer is a parameter learned while training. A convolutional layer in our network therefore has $\text{input dimension} \times \text{filter width} \times \text{filter height} \times \text{number filters}$ parameters. Our LeNet-like network with 16 filters per convolutional layer has about 5.000 parameters in the convolutional layers and about 3.500 ones in the fully connected parts. By exponential increase in the number of filters we exponentially increase the number of parameters and therefore the training complexity. The dominance of parameters in the convolutional layers versus the fully connected layer increases dramatically.

While experimenting with various filter sizes I had especially problem with the limited memory size on GPU. While prediction e.g. for classification error calculation, the full MNIST training dataset got processed. 3GB gRAM was too small for the high dimensional output of convolution layers with high number of filter. As a workaround I modified the code to always process batches of the dataset.

In Figure 2 you'll see how different number of filters affects the training time on CPU and GPU. The CPU is always slower than the GPU. We expect this because the GPU is optimized for matrix operations and parallel executions. This is handy for forward and backpropagation of our network. For the exponential increase in filter number the CPU runtime increases also exponential-like. The CPU has to process the operations sequentially wherefore double the time is necessary for double the parameters. On GPU runtime just slightly increases by increasing filter number up to 2^5 . The GPU can compensate the increase of parameters by massive parallel processing. Still, for higher filter numbers all parallel processing units are in use such that the additional ones have to be processed in sequence as well. Therefore runtime ascends quickly for higher number of filters.




Figure 2: Runtime for 10 training epochs with LeNet-like CNN and MNIST dataset on GPU and CPU. The number of filters per convolutional layer was modulated. The code was executed on a workstation with *Intel®Core™i7-3770 CPU @ 3.40GHz* and *NVIDIA GeForce GTX 1060*.